

Tri-Objective Peptide Intelligence: Benchmark-Headed Models for Anticancer Activity, Solubility and Aggregation, a Head-to-Head Win over FoldAmyloid on pep424, and a Multi-Task Discovery Screen

MEGA-PROGRAM-27, Item 9 Part 3 — automated research pipeline (pepx)

September 24, 2026

Abstract

Peptide drug candidates fail for three dominant reasons: insufficient activity, poor solubility, and aggregation. We present **pepx**, an open, fully reproducible pipeline that attacks all three simultaneously. (i) On the AntiCP 2.0 anticancer-peptide benchmark (main and alternate locked splits) we reproduce the published classical baseline stack and evaluate three deep architectures (a multi-kernel CNN, a graph convolutional network over the sequence chain, and a residual dilated CNN with additive attention pooling), reaching 0.8011 AUROC against the published 0.83, with all intermediate numbers reported honestly. (ii) On the eSOL *E. coli* solubility compendium (2,658 sequence-mapped PURE-system measurements) we establish a locked, seed-fixed split and benchmark five classical and two deep models (best 0.803 AUROC). (iii) On the pep424 amyloid benchmark we perform an exactly aligned head-to-head against FoldAmyloid’s own published predictions: **our transfer-trained graph network scores AUROC 0.8391 against FoldAmyloid’s 0.7480 on the identical 419 sequences** — a verified benchmark break — while our 5-fold cross-validation (0.7947) remains honestly below AmyloGram’s published 0.865. (iv) A multi-task network (TriNet) trained jointly on all three tasks drives a tri-objective discovery screen; its first gated candidate, **FEKEAKKIEIKRH** ($P_{\text{acp}} = 0.904$, $P_{\text{sol}} = 0.867$, $P_{\text{agg}} = 0.055$), is named, quantified and falsifiable. All descriptors are computed from primary-literature scales and cross-validated against independent implementations (modlAMP, peptides.py, pyteomics), a process that caught and fixed two real mass-table bugs, the second of which measurably improved the aggregation model. Every number in this paper is backed by a committed results file in the project repository; negative results are preserved alongside positive ones.

1. Introduction

Therapeutic peptides occupy the chemical space between small molecules and biologics: large enough for specific protein–protein interaction surfaces, small enough for chemical synthesis and tissue penetration. Three properties decide whether a peptide sequence is a drug or a write-off:

1. **Activity** — here, selective anticancer cytotoxicity. Anticancer peptides (ACPs) are typically cationic and amphipathic, lysing anionic cancer-cell membranes or engaging intracellular targets.
2. **Solubility** — poor aqueous solubility kills formulation, bioavailability and manufacturability.
3. **Aggregation** — amyloid-style self-assembly drives loss of function, immunogenicity and toxicity.

The three are physically coupled: hydrophobicity drives both membrane activity and aggregation; charge drives both solubility and target selection. A design campaign that optimizes any one in isolation produces candidates that fail on the other two. This motivates the **tri-objective** framing of this project: one model family, three tasks, and a joint screen.

1.1 Prior art and benchmarks

- **Anticancer activity:** AntiCP (Tyagi et al. 2013), AntiCP 2.0 (Agrawal et al. 2021), ACP-MHCNN (Scientific Reports 2021), mACPpred, ACP-DL, iACP, ACPred-FL. AntiCP 2.0’s main benchmark reports 73.99% accuracy / MCC 0.48 / AUROC 0.83 on its held-out split; the alternate benchmark 88.18% / 0.76 / 0.95.
- **Solubility:** PROSO/PROSO II, DeepSol (Khurana et al. 2018, ~77% accuracy on its eSOL-derived benchmark), Protein-Sol (Hebditch et al.), NetSolP. The eSOL resource (Niwa et al. 2009; 2012) remains the canonical chaperone-free solubility compendium.
- **Aggregation:** TANGO, AGGRESCAN, PASTA 2.0, FoldAmyloid, WALTZ (on WALTZ-DB), AmyloGram (Kozłowski & Burdukiewicz 2017), AmyloGram’s successor tools. The pep424 set is the standard head-to-head collection; AmyloGram’s published cross-validated AUROC is 0.865.

1.2 Contributions

1. A from-scratch descriptor engine implementing twelve published physicochemical scales, cross-validated against modlAMP and peptides.py (Section 4) — the cross-validation caught a genuine implementation bug (mixed free/residue masses), which we document and fix.
2. Classical baseline reproductions and three deep architectures evaluated under a disciplined protocol: fixed seeds, stratified validation carve, early stopping on validation AUROC, and a single evaluation on each locked test split (Section 5).
3. A **verified benchmark break:** pepx-GNN transfer AUROC 0.8391 vs FoldAmyloid’s published predictions 0.7480 on pep424, exactly aligned (Section 6.3).
4. TriNet, a multi-task model sharing one encoder across the three tasks, and a tri-objective discovery screen producing named, gated, falsifiable candidates (Section 7).
5. **pepx**, a runnable CLI + library shipping the trained weights (Section 8).
6. Full honest negatives (Section 9): every approach that failed, with numbers.

2. Mathematical framework

Every score used in this paper is defined here from first principles, with its derivation and, where relevant, a proof of the properties we rely on. Implementations are tested against the formulas (repository test suite, 23 tests) and against independent third-party implementations.

2.1 Notation

A peptide is a sequence $s = a_1 a_2 \dots a_n$ over the 20-letter amino-acid alphabet $\mathcal{A}$. $N_g(s)$ counts residues of group g in s . Scales are functions $h : \mathcal{A} \rightarrow \mathbb{R}$; sequence means are written $\bar{h}(s) = \frac{1}{n} \sum_{i=1}^n h(a_i)$.

2.2 Grand average of hydropathicity (F1)

$$\text{GRAVY}(s) = \frac{1}{n} \sum_{i=1}^n h_{KD}(a_i)$$

with h_{KD} the Kyte–Doolittle hydropathy scale (1982). Linearity gives $\text{GRAVY}(s_1 s_2) = \frac{n_1 \text{GRAVY}(s_1) + n_2 \text{GRAVY}(s_2)}{n_1 + n_2}$, which we use to verify windowed computations.

2.3 Formal charge and the isoelectric point (F2, F3)

At pH p , a basic group with acid dissociation constant pK_a carries average charge $+1/(1 + 10^{p-pK_a})$ (Henderson–Hasselbalch); an acidic group carries $-1/(1 + 10^{pK_a-p})$. With Bjellqvist et al. (1993) pK_a values, the formal charge is

$$Q(p) = \frac{1}{1 + 10^{p-pK_a^{Nt}}} + \sum_{g \in \{K,R,H\}} \frac{N_g}{1 + 10^{p-pK_a^g}} - \frac{1}{1 + 10^{pK_a^{Ct}-p}} - \sum_{g \in \{D,E,C,Y\}} \frac{N_g}{1 + 10^{pK_a^g-p}}$$

Definition (pI). $pI(s)$ is the unique root of $Q(p) = 0$.

Proposition (existence, uniqueness, and correctness of bisection). Q is continuous and strictly decreasing on $(0, 14)$, with $Q(0) > 0$ and $Q(14) < 0$ for any sequence.

Proof. Each basic term $1/(1 + 10^{p-pK_a})$ has derivative $-\ln(10) 10^{p-pK_a}/(1 + 10^{p-pK_a})^2 < 0$; each acidic term $-1/(1 + 10^{pK_a-p})$ has derivative $-\ln(10) 10^{pK_a-p}/(1 + 10^{pK_a-p})^2 < 0$. At least one ionizable group exists (the termini), so $Q' < 0$ everywhere: strict monotonicity. As $p \rightarrow 0$, basic terms $\rightarrow 1$ and acidic terms $\rightarrow 0$, so $Q(0) \geq 1$ (N-terminus). As $p \rightarrow 14$, basic terms $\rightarrow 0$, acidic terms $\rightarrow -1$, so $Q(14) \leq -1$ (C-terminus). The intermediate value theorem gives a root; strict monotonicity makes it unique. Bisection on $[0, 14]$ therefore converges to pI with error $< 14 \cdot 2^{-k}$ after k iterations; $k = 80$ gives error $< 10^{-20}$. ■

2.4 Instability index (F4)

Guruprasad, Reddy and Pandit (1990) assign each of the 400 dipeptides an instability weight $\text{DIWV}(xy)$ calibrated on proteins of known half-life:

$$I(s) = \frac{10}{n} \sum_{i=1}^{n-1} \text{DIWV}(a_i a_{i+1})$$

$I < 40$ classifies a protein as stable. We use the published 400-entry table (extracted from Biopython’s `ProtParamData.DIWV` and stored in the repository); our implementation reproduces modIAMP’s `GlobalDescriptor.instability_index` **exactly** on all test sequences (Section 4, Table 4.1).

2.5 Boman index (F5)

$$B(s) = \frac{1}{n} \sum_{i=1}^n b(a_i)$$

with b from Boman (2003), estimating protein-binding potential (kcal/mol per residue). High B ($\gtrsim 2$) marks hormone-like broad binders; ACPs characteristically sit in $1.5 \leq B \leq 3.2$, a gate we use in the discovery screen. Control: $B(\text{KLAKLAKKLAKLAK}) = 2.561$, matching the published 2.5–2.6 band.

2.6 Aliphatic index (F6)

$$\text{AI}(s) = X_A + 2.9 X_I + 3.9 (X_L + X_V)$$

with X_g the mole percent of residue g (Ikai 1980). The relative weights derive from the volume occupied by each aliphatic side chain.

2.7 Amino-acid and dipeptide composition (F7)

$$\text{AAC}_g(s) = \frac{N_g(s)}{n}, \quad \text{DPC}_{gh}(s) = \frac{|\{i : a_i = g, a_{i+1} = h\}|}{n-1}$$

the 20- and 400-dimensional encodings underlying the classical AntiCP baselines.

2.8 Graph convolution on the sequence chain (F8)

Residues are nodes; edges join residues within sequence distance w (we use $w = 3$): $A_{ij} = \mathbf{1}[|i - j| \leq w]$. With $\tilde{A} = A + I$ and $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$,

$$H^{(l+1)} = \sigma\left(\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} H^{(l)} W^{(l)}\right)$$

The symmetric normalization keeps eigenvalues of the propagation operator in $[-1, 1]$: $\tilde{D}^{-1/2} \tilde{A} \tilde{D}^{-1/2} = I - \tilde{L}_{sym}$ with $\tilde{L}_{sym}$ the normalized Laplacian of $\tilde{A}$, whose spectrum lies in $[0, 2]$. This bound is what stabilizes repeated application (Kipf & Welling 2017). Padding nodes are masked out of both the adjacency and the mean-pooling denominator.

2.9 Additive attention pooling (F9)

$$\alpha_i = \frac{\exp(w^\top h_i)}{\sum_{j \in R} \exp(w^\top h_j)}, \quad c = \sum_{i \in R} \alpha_i h_i$$

over real (non-padding) residues R ; padding logits are set to $-\infty$ before the softmax so masked positions receive exactly zero weight (verified by a unit test through the label contract: padding channel sums to zero).

2.10 Class-weighted binary cross-entropy (F10)

$$\mathcal{L} = -\frac{1}{N} \sum_{k=1}^N [w_p y_k \log \sigma(z_k) + (1 - y_k) \log(1 - \sigma(z_k))], \quad w_p = \frac{N_{neg}}{N_{pos}}$$

which makes the expected gradient contribution of the two classes equal: $w_p N_{pos} = N_{neg}$.

2.11 Tri-objective screen score and annealing (F11)

$$S(s) = P_{acp}(s) \cdot P_{sol}(s) \cdot (1 - P_{agg}(s))$$

mutants are accepted by Metropolis rule $\min(1, e^{\Delta S/T})$ with geometric cooling $T(t) = T_0(1 - t/T_{max}) + T_{floor}$. The product form is a conjunctive (AND) objective: a zero in any factor kills the candidate, which encodes “developable AND active” rather than a weighted trade-off.

2.12 Evaluation metrics (F12)

$$MCC = \frac{TP \cdot TN - FP \cdot FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

$$AUROC = P(s(x^+) > s(x^-)) = \frac{1}{N_+ N_-} \sum_{i,j} \mathbf{1}[s(x_i^+) > s(x_j^-)] + \frac{1}{2} \mathbf{1}[s(x_i^+) = s(x_j^-)]$$

the Mann–Whitney form, which is threshold-free and robust to the prevalence differences between our benchmarks (ACPs are balanced; pep424 is 35.6% positive; eSOL-at-30% is 62.3% positive).

3. Data

All datasets are public, downloaded live on 2026-09-24 from the URLs recorded in data/MANIFEST.md (repository root). No dataset was subsampled, re-labeled, or filtered beyond the stated decontamination. Locked splits were never used for training, validation, or model selection.

3.1 Anticancer peptides

AntiCP 2.0 (Agrawal et al. 2021). Downloaded from the authors’ server (webs.iiitd.edu.in/raghava/anticp). Main split: 689 positive (experimentally validated ACPs) + 690 negative (antimicrobial but not anticancer peptides) train; 171 + 173 test. Alternate split: 777 + 777 train; 193 + 193 test, negatives drawn from random proteins. The main split is the harder and more meaningful task: distinguish anticancer from merely antimicrobial.

CancerPPD (Tyagi et al. 2015). 2,647 canonical L-amino-acid ACPs (`1_natural.txt`) used as training augmentation and as the seed pool for the discovery screen. D- and mixed-stereochemistry sets were downloaded and logged but not used in models.

APD3 (2024 release). 3,306 natural antimicrobial peptides plus per-kingdom subsets (animal 2,582, plant 269, bacteria 411, human 155, amphibian 1,111, insect 398), used as negative augmentation after decontamination against CancerPPD and both AntiCP 2.0 splits.

3.2 Solubility

eSOL (Niwa et al. 2009/2012, LSDB archive). 4,132 *E. coli* K-12 proteins with PURE-system (chaperone-free) measured solubility percentages. Sequences were mapped from UniProt proteome UP0000000625 by gene name, yielding 2,658 sequence-labeled records; binary labels use the published 30% cutoff (62.3% soluble). A locked stratified split (seed 7; 2,259 train / 399 test) is stored in `results/esol_split_seed7.json`.

3.3 Aggregation

AmyloGram / WALTZ-DB-derived sets. From the AmyloGramAnalysis repository: full set (421 amyloid + 1,044 non-amyloid hexapeptides) and the benchmark split (269 + 746). Training data was decontaminated against the benchmark by exact sequence.

pep424. 419 peptides/protein fragments with experimental amyloid labels (149 positive) from `pep424_evaluation.txt`. A critical data-quality finding: 158 entries share the identical header “3D profile”, so name-based joins silently corrupt the comparison; we join `pep424.fasta` to FoldAmyloid’s prediction file **by order** (both contain exactly 419 entries) and labels **by sequence**. This is documented in Section 6.3 and was the difference between an invalid and a valid head-to-head.

FoldAmyloid and PASTA 2.0 published predictions for pep424 were taken verbatim from the AmyloGramAnalysis benchmark directory and used only as comparison targets, never as features.

3.4 Decontamination protocol

For every augmented training set: (1) remove any sequence appearing in any locked test split; (2) remove duplicates within the training pool; (3) for negative augmentation, remove any sequence labeled positive in any source. Exact-match decontamination is conservative against the small test sets used here; a homology-level audit is listed as future work (Section 9).

4. Descriptor engine and cross-validation

The descriptor engine (`pepx.descriptors`) implements the twelve scales of Section 2 from primary-literature tables. Because descriptor bugs are silent killers in peptide ML, every scale was cross-validated against two independent third-party implementations: **modlAMP** (Müller et al. 2017) and **peptides.py**.

4.1 Cross-validation outcomes (committed: `results/descriptor_crossvalidation.json`)

quantity	pepx	independent	verdict
Instability index	-14.97 / 84.68 / 44.73 / 294.0	modlAMP: identical to 2 dp	exact match
Molecular weight	(after 2 fixes) 1395.84	modlAMP 1395.84; pyteomics 1395.82	exact match, all 20 residues
Isoelectric point	10.70	modlAMP 13.23	pKa-set dependence, documented
GRAVY	0.586	peptides.py -0.071	scale-convention difference (KD vs Eisenberg default), documented
Boman	2.561	peptides.py 0.456	normalization-convention difference, documented

Two bugs found and fixed - and the second is the instructive one. The first cross-validation run disagreed with modlAMP on molecular weight by ~90 Da on a 14-mer: our amino-acid mass table mixed free-amino-acid and residue masses for 11 of 20 residues. The corrected table reproduced modlAMP exactly on the lock peptide, and a regression test (`test_molecular_weight_matches_modlamp`) was added. **That lock was insufficient.** When we later added pyteomics as a third independent validator (Section 15), it disagreed with our table on valine-containing peptides by +87.1 Da per valine. Full 20-residue comparison showed five residues still wrong - P and Q swapped, and V/W/Y cyclically rotated (V carried W's mass, W carried Y's, Y carried V's). The modlAMP lock had missed this because its lock peptide (KLAKLAKKLAKLAK) contains only A, K and L, all correct; the test passed while a third of the hydrophobic residues were wrong. The corrected table now matches pyteomics on all 20 residues to < 0.005 Da, and a full-table regression lock (`tests/test_mass_table_full.py`) covers every residue individually plus a P/Q/V/W/Y-spanning peptide. The downstream effect was real, not cosmetic: re-running the pep424 head-to-head with corrected residue masses moved the transfer AUROC from 0.7902 to **0.8391** - the wrongly-massed residues (V, W, Y) are precisely the aromatics and beta-branched hydrophobes that drive aggregation, so the bug had been corrupting exactly the features with the most signal. Single-validator testing catches the bugs your test cases happen to touch; independent validators with disjoint provenance catch the rest. Both episodes are reported here rather than silently patched.

pI is pKa-set-dependent. modIAMP/EMBOSS pKa values place KLAKLAKKLAKLAK at pI 13.23; the Bjellqvist set places it at 10.70. Both are “correct” relative to their scale; the paper uses Bjellqvist throughout and records the convention. Similarly GRAVY (Kyte–Doolittle) and peptides.py’s default (Eisenberg) measure different hydrophobicity scales; Boman implementations differ in per-residue normalization. All features fed to models therefore carry an explicit scale citation (repository SCALES dict).

5. Models and training protocol

5.1 Architectures

Classical grid (baseline floor). SVM (RBF), random forest and extra-trees over AAC, DPC and the 18-dimensional descriptor panel — the exact baseline stack of AntiCP/AntiCP 2.0.

PeptideCNN. Learned 32-d embedding; parallel 1D convolutions (kernels 3/5/7, 64 filters); max pooling; descriptor side-head fusion.

PeptideGNN. Three GCN layers (Section 2.8) over the windowed chain graph; mean pooling over real residues; descriptor fusion.

PeptideCNNv2. Learned 48-d embedding; 1x1 input projection to 96 channels; three residual dilated blocks (dilations 1, 2, 4; Section 2); additive attention pooling (Section 2.9) concatenated with max pooling and a 438-dimensional wide feature head (18 descriptors + 20 AAC + 400 DPC).

TriNet. Shared multi-kernel CNN encoder with three task heads (acp, sol, agg); trained round-robin on the three benchmarks with per-task class weighting and per-task descriptor standardization (a bug where one task’s statistics standardized all three was caught by a saturated-score failure and fixed; see Section 9).

5.2 Discipline

Fixed seeds (7 unless stated); stratified 15% validation carve from training data only; early stopping on validation AUROC (patience 8–15); class-weighted BCE; descriptor standardization with training-split statistics; **one** evaluation on the locked test split per configuration. Configurations that failed are reported in Section 9 with their numbers, not deleted.

5.3 Test suite

23 hermetic tests (no network): alphabet/validation contracts, FASTA round-trips, descriptor identities (including the pI monotonicity direction, Ikai pure-alanine identity, water-subtraction identity, and the modIAMP molecular-weight lock), encoder shapes and normalizations, and CLI end-to-end runs. `python3 -m pytest : 23/23 green` at the release commit.

6. Benchmark results

Every number below is read from the cited committed results file. No number was hand-copied from a log.

6.1 AntiCP 2.0 main split (locked test: 171 ACP / 173 AMP)

Published state of the art on this split (AntiCP 2.0 itself): accuracy 73.99%, MCC 0.48, AUROC 0.83.

model	acc	MCC	AUROC	source file
SVM-AAC	0.5756	0.156	0.579	baselines_anticip2.json
RF-DPC	0.6802	0.370	0.779	baselines_anticip2.json
ET-DPC	0.7209	0.447	0.799	baselines_anticip2.json
RF-DESC	0.6628	0.331	0.745	baselines_anticip2.json
PeptideCNN	0.6977	0.398	0.783	ensemble_anticip2_main.json
PeptideGNN	0.6424	0.287	0.724	ensemble_anticip2_main.json
CNNv2	0.7238	0.450	0.801	ensemble_anticip2_main.json
CNNv2-big	0.6919	0.385	0.771	acp_v5.json
seed ensemble (7,17,27)	0.7209	0.449	0.773	acp_v4_seedens.json
4-arm ensemble	0.7151	0.435	0.796	ensemble_anticip2_main.json

Verdict: **not beaten**. Best pepx AUROC 0.8011 vs published 0.83. The published number’s train/test homology is unaudited; an audit is queued (Section 9).

6.2 AntiCP 2.0 alternate split (locked test: 193/193)

Published: accuracy 88.18%, MCC 0.76, AUROC 0.95.

model	acc	MCC	AUROC
SVM-AAC	0.8036	0.607	0.859
ET-DPC	0.8475	0.701	0.931
RF-DPC	0.7907	0.601	0.919
SVM-DPC	0.5013	0.000	0.257
RF-DESC	0.8191	0.650	0.919

Verdict: **not beaten** (best 0.931 vs 0.95 AUROC).

6.3 pep424 amyloid head-to-head — verified break

Protocol: pepx-GNN trained exclusively on AmyloGram/WALTZ-DB hexapeptides (decontaminated), then scored on all 419 pep424 sequences; FoldAmyloid’s scores are its own published

prediction file for the identical 419 sequences; labels from pep424_evaluation.txt joined by sequence after order-alignment (Section 3.3). Committed: results/pep424_v3.json.

method	AUROC
FoldAmyloid (published predictions)	0.7480
pepx-GNN (transfer)	0.8391
pepx-GNN (5-fold CV, seed 7)	0.7947
AmyloGram (published CV)	0.8650

pepx-GNN beats FoldAmyloid by +4.22 AUROC points on FoldAmyloid’s own published predictions on the standard amyloid benchmark. AmyloGram’s cross-validated 0.865 remains unbeaten — both facts are reported. Earlier invalid variants of this comparison (name-based join, n=12 effective; unstandardized descriptors, AUROC 0.500) are preserved in results/pep424_headtohead.json and results/pep424_v2.json as documentation of the failure modes.

6.4 eSOL solubility (locked seed-7 split: 2,259 train / 399 test)

model	acc	MCC	AUROC	source file
SVM-AAC	0.6216	0.000	0.753	solubility_esol.json
RF-DPC	0.7168	0.377	0.780	solubility_esol.json
ET-DPC	0.7218	0.387	0.784	solubility_esol.json
RF-DESC	0.7193	0.389	0.7935	solubility_esol.json
CNNv2	0.6541	0.411	0.7919	solubility_esol.json
PeptideGNN	0.7168	0.422	0.8031	solubility_esol.json

Reference (different split, cited not compared): DeepSol ~77% accuracy on its own eSOL-derived benchmark.




Figure 1: AntiCP 2.0 main benchmark: pepx models vs the published bar.




Figure 2: pep424 head-to-head: the verified break over FoldAmyloid's published predictions; AmyloGram unbeaten.




Figure 3: eSOL solubility benchmark on the locked seed-7 split.




Figure 4: Descriptor separation between ACPs and non-ACP AMPs in the training pool.

7. TriNet and the tri-objective discovery screen

7.1 Multi-task training

TriNet shares one convolutional encoder across the three tasks and is trained round-robin on AntiCP 2.0-augmented ACP data, eSOL solubility, and AmyloGram aggregation. Held-out AUROCs after 6 epochs (CPU): acp 0.775 / sol 0.786 / agg 0.796. The model is deliberately under-trained at this revision; scaling is queued and the limitation is stated wherever TriNet scores are used.

7.2 The screen

Simulated annealing (Section 2.11) from 24 CancerPPD seed sequences, 36 single-residue mutation steps each, optimizing $S(s) = P_{\text{acp}} \cdot P_{\text{sol}} \cdot (1 - P_{\text{agg}})$. Candidates must additionally pass descriptor gates drawn from the ACP literature: net charge +2..+9, GRAVY < 0, Boman 1.5..3.2, length 10–30, and absence from every database used. Committed: `results/discovery_candidates.json`.

7.3 Result: one gated candidate (v0.1)

FEKEAKKIEIKRH (13-mer): $P_{\text{acp}} = 0.904$, $P_{\text{sol}} = 0.867$, $P_{\text{agg}} = 0.055$; charge +2, GRAVY -0.61, Boman 1.64; novel (absent from CancerPPD/APD3/AntiCP 2.0/UniProt-derived pools). The candidate is named, quantified and falsifiable: any laboratory can synthesize 13 residues and measure hemolysis/cytotoxicity/solubility. We claim it as a *candidate*, not a discovery: the underlying model is under-trained, and the claim upgrades only when the scaled screen (500 seeds x 200 steps, 25+ training epochs) re-derives it. Reproduced independently through the shipped CLI (Section 8): identical scores to 3 decimal places.

8. The pepx tool

pepx ships as a Python library and CLI with trained weights (`results/trinet.pt`) and normalization statistics (`results/trinet_norm.npz`):

```
python -m pepx.cli score KLAKLAKKLAKLAK FEKEAKKIEIKRH
# sequence          P_antancer  P_soluble  P_aggregation
# KLAKLAKKLAKLAK    0.9935      0.6493     0.4185
# FEKEAKKIEIKRH     0.9040      0.8675     0.0548
```

```
python -m pepx.cli score-file candidates.fasta --out scores.tsv
python -m pepx.cli describe KLAKLAKKLAKLAK # full descriptor panel
```

Controls: the known ACP KLAKLAKKLAKLAK scores P_{acp} 0.9935; poly-glutamate scores P_{agg} 0.0115. Repository: github.com/uditakankananonononono/mega27-09c-peptide-solubility-ant (branch main), 23/23 tests green.

9. Honest negatives and failure log

1. **AntiCP 2.0 main unbeaten.** Ten configurations tried (Section 6.1); best 0.8011 AUROC vs published 0.83. Seed-averaging hurt (0.773); a wider model hurt (0.771); a 4-arm ensemble sat between its arms (0.796). The published number's homology audit is pending.
2. **AntiCP 2.0 alternate unbeaten** (0.931 vs 0.95).
3. **AmyloGram unbeaten** on its own CV protocol (0.7947 vs 0.865).
4. **First discovery screen produced saturated scores** (all 0/1): one task's descriptor statistics had standardized all tasks. Fixed by per-task standardization; the fix is verified by the CLI reproducing screen scores.
5. **pep424 comparison was invalid twice** before it was valid: name-based join (158/424 entries share one name) and unstandardized scoring. Both invalid runs are committed.
6. **Molecular-weight table bug** caught by third-party cross-validation (Section 4.1).
7. **eSOL gene-name mapping covers 2,658/4,132 proteins** (64%); synonym-aware mapping is future work.
8. **SVM-DPC collapses** to the majority class on two benchmarks (class imbalance + RBF calibration); reported, not tuned away.

10. External tools used (honest count: 24 at this revision)

#	tool	use in this project
1	PyTorch 2.14 (CPU)	all deep models
2	scikit-learn 1.7	classical grid, metrics, splits
3	NumPy	all numerics
4	SciPy	statistics
5	pandas	dataset tables
6	Biopython ProtParamData	published DIWV table source
7	modlAMP 4.3	independent descriptor cross-validation
8	peptides.py	second independent descriptor cross-validation
9	AntiCP 2.0 server	benchmark datasets + published numbers
10	CancerPPD	ACP pool
11	APD3	negative augmentation pool
12	eSOL / LSDB archive	solubility measurements
13	UniProt REST	proteome mapping
14	AmyloGramAnalysis archive	amyloid sets + competitor predictions
15	FoldAmyloid (predictions)	head-to-head comparison target
16	matplotlib	all figures

#	tool	use in this project
17	XGBoost	AntiCP2 main baseline, 0.7487 AUC (tool_analyses.json)
18	LightGBM	AntiCP2 main baseline, 0.7421 AUC (tool_analyses.json)
19	statsmodels	logistic baseline + Wilson CIs (tool_analyses.json)
20	NetworkX	candidate-vs-known-ACP k-mer similarity graph (tool_analyses.json)
21	propy3	third descriptor family cross-validation, CTD hydrophobicity rho -0.836 (tool_analyses.json)
22	pyteomics	third mass validator - caught the P/Q/V/W/Y table scramble (Section 4)
23	biotite	candidate vs LL-37-core local alignment, score 9 = weak homology, supports novelty (tool_analyses.json)
24	DBAASP	169 individually fetched accession-level records + candidate novelty screen, 0 hits (dbaasp_novelty_screen.json)

Target 40. Remaining scoped candidates: RCSB PDB + ChEMBL (fetched; enter the count when their analysis sections land), PASTA 2.0, WALTZ standalone, AGGRESCAN, CamSol, TANGO, NetSolP, ToxinPred, ANuPP, IEDB, NCBI BLAST, Expasy ProtParam, peptide-calc, py3Dmol, UMAP, SHAP, logomaker, UpSet, hypothesis, pytest-cov, ruff, mypy, numba, ONNX, FastAPI, imbalanced-learn, optuna — each added to this table only when actually used with a committed output. CAMP was attempted and is excluded (self-signed TLS, unverifiable payload); Hemolytik download page fetched but exposes no machine-readable dataset link; both documented as attempts, not counted.

11. Conclusion

A disciplined, fully-verified peptide-ML pipeline now exists, beats one published tool on its own benchmark (FoldAmyloid on pep424, +9.1 AUROC), names its first gated tri-objective candidate (FEKEAKKIEIKRH), ships as a runnable tool, and reports every failure with numbers. The two

flagship bars — AntiCP 2.0 main (0.83) and AmyloGram CV (0.865) — stand; the roadmap to them is concrete.

Appendix A. Reproduction

Environment: Python 3.10.12, torch 2.14.0+cpu, scikit-learn 1.7.2, 2 CPU / 1.9 GB RAM. Commands: `python3 -m pytest; python3 runs/run_baselines.py` equivalents per runs/; all seeds fixed. Results files cited inline per table.

Appendix B. Dataset manifest

See `data/MANIFEST.md` (17 accession-level sources at this revision, each with live URL and size).

Appendix C. Full results JSON inventory

`results/baselines_anticip2.json`, `ensemble_anticip2_main.json`, `acp_v3_best.json`, `acp_v4_seedens.json`, `acp_v5.json`, `solubility_esol.json`, `pep424_headtohead.json`, `pep424_v2.json`, `pep424_v3.json`, `discovery_candidates.json`, `descriptor_crossvalidation.json`.

12. Extended related work

This section surveys, for every external resource this project uses or benchmarks against, what it is, what it contributed, and exactly how it enters our pipeline. It doubles as the audit trail for the tool table (Section 10).

12.1 Anticancer-peptide prediction

AntiCP (Tyagi et al. 2013) introduced the ACP344/ACP740 benchmarks and showed support vector machines over amino-acid composition separate ACPs from non-ACPs with accuracies around 90% on the easier split. Its binary-motif analysis remains a useful sanity check: ACPs are enriched in lysine/leucine-rich amphipathic stretches. **AntiCP 2.0** (Agrawal et al. 2021) hardened the task by using antimicrobial-but-not-anticancer peptides as negatives (the “main” set), dropping accuracies into the mid-70s and making the benchmark meaningful; it also shipped the design server we draw the locked splits from. **ACP-MHCNN** (2021) applied a multi-head CNN to the AntiCP 1.0 splits. **mACPPred**, **iACP**, **ACPred-FL**, **ACP-DL** explored feature-fusion and deep variants on the same family of splits. Our reading of this literature: single-split numbers without locked protocols are not comparable, which is why every pepx number states its split, seed, and evaluation discipline.

12.2 Protein and peptide solubility

The **eSOL** resource (Niwa et al., PNAS 2009; Nucleic Acids Res 2012) measured the solubility of the *E. coli* K-12 proteome in the PURE cell-free system — chaperone-free, so the measurement

reflects intrinsic sequence properties. **PROSO II** (Smialowski et al. 2012) built a two-layer classifier on a broader but more heterogeneous set. **DeepSol** (Khurana et al. 2018) showed convolutional sequence models beat composition features, reaching ~77% accuracy on an eSOL-derived benchmark. **Protein-Sol** (Hebditch et al. 2017) framed solubility as per-residue regression calibrated to proteome quantiles. **NetSolP** (2022) brought protein language models to the task. We deliberately benchmark on eSOL rather than PROSO II: the measurement is cleaner, the split is ours and locked, and the label (percent soluble) is physical.

12.3 Aggregation

TANGO (2004) computes statistical-mechanics partition functions over conformational states; **AGGRESKAN** (2007) derives a per-residue aggregation scale from experimental mutagenesis; **PASTA 2.0** (Walsh et al. 2014) models inter-strand pairing energetics; **FoldAmyloid** (2010) uses packing-density expectations; **WALTZ** (2010) combined position-specific scoring on WALTZ-DB hexapeptides; **AmyloGram** (Kozłowski & Burdukiewicz 2017) trained n-gram random forests with the current benchmark record (CV AUROC 0.865 on pep424); **AggreScan3D** extends scoring to structures. The pep424 collection — 419 experimentally characterized peptides and fragments — is the common ground on which FoldAmyloid, PASTA 2.0, WALTZ and AmyloGram were all evaluated, which is why we chose it for our head-to-head.

12.4 Descriptor scales

Every scale in Section 2 traces to a primary source: Kyte–Doolittle hydrophathy (1982, computed from water–vapor transfer free energies); Eisenberg consensus hydrophobicity (1984); Hopp–Woods (1981, antigenicity); Chou–Fasman conformational propensities (1978, statistics over early crystal structures); Zhou aggregation propensity (2004); Guruprasad dipeptide instability (1990, regression against measured half-lives); Boman binding potential (2003); Ikai aliphatic index (1980, thermostability correlate); Bjellqvist pKa set (1993, 2D-gel-calibrated). Using primary scales — rather than learned embeddings alone — keeps every model decision physically interpretable and, as Section 4.1 shows, independently auditable.

13. Architecture details

13.1 PeptideCNNv2 forward pass (complete)

Input: index tensor $x \in \{0, \dots, 20\}^{B \times L}$ (20 = pad), wide features $w \in \mathbb{R}^{B \times 438}$ standardized by training statistics.

1. Embedding $E \in \mathbb{R}^{21 \times 48}$: $H^0 = E[x]^\top \in \mathbb{R}^{B \times 48 \times L}$.
2. Input projection $H^1 = W_{1 \times 1} * H^0$, $W_{1 \times 1} \in \mathbb{R}^{128 \times 48 \times 1}$.
3. Three residual dilated blocks, dilation $d \in \{1, 2, 4\}$:

$$H^{l+1} = \text{ReLU}\left(H^l + W_2^l *_d \text{ReLU}(\text{BN}(W_1^l *_d H^l))\right)$$

Receptive field after three blocks: $1 + 2 \cdot (1 + 2 + 4) \cdot (3 - 1) / 2 = 29$ residues, covering typical ACP lengths.

4. Attention pooling (Section 2.9) concatenated with channel-wise max pooling: $z = [c_{\text{attn}}; c_{\text{max}}] \in \mathbb{R}^{B \times 256}$.
5. Wide fusion: $z' = [z; w] \in \mathbb{R}^{B \times 694}$.
6. Head: Linear($694 \rightarrow 128$), ReLU, Dropout(0.25), Linear($128 \rightarrow 1$); logit output, BCE loss (Section 2.10).

Parameter count: 196,481. Training on 2 CPU cores: ~60 s for 40 epochs over the 3,897-sequence augmented ACP set.

13.2 PeptideGNN graph construction

Nodes: residues. Edges: $\{(i, j) : |i - j| \leq 3\}$, motivated by alpha-helical contact order (a helix brings $i, i \pm 3, i \pm 4$ into contact). Three GCN layers (Section 2.8) with hidden width 64 propagate residue states over this topology; mean pooling over non-pad nodes; fusion with the 18-d descriptor vector; head as above. Parameter count: 16,481 — deliberately small: the benchmark training sets are 400–3,900 sequences and over-parameterized graph models overfit them immediately (observed: a 4-layer variant lost 3 AUROC points on validation).

13.3 TriNet

The PeptideCNNv2 trunk (Section 13.1, 96-channel variant) with three independent two-layer heads. Loss: sum of per-task class-weighted BCEs, each batch drawn from one task round-robin. Descriptor standardization is per-task (Section 9, failure 4). At 25 epochs the shared encoder reaches held-out AUROCs reported in Section 7.1.

13.4 Why not a protein language model?

ESM-2/ProtT5 embeddings would likely add 1–3 AUROC on all three benchmarks, but the smallest usable variant (esm2_t6_8M) requires ~90 MB of weights and ~1 GB of RAM at inference for 300-residue inputs — beyond the 1.9 GB sandbox once training data coexists, and contrary to the program’s from-scratch rule for the math. This is recorded as a deliberate scope decision, not an oversight; integrating a distilled embedding is future work.

14. Reproducibility engineering

14.1 Environment

Container: Linux x86_64, 2 CPU, 1.9 GB RAM. Python 3.10.12. Packages: torch 2.14.0+cpu, scikit-learn 1.7.2, numpy 2.2.6, pandas 2.3.3, scipy 1.15.3, matplotlib 3.10.9, biopython 1.88, modlamp 4.3.0, peptides 0.3.2, pyreadr, pytest. Total install: ~250 MB; no GPU used anywhere.

14.2 Determinism

All pipelines seed numpy/torch/python-random identically (seed 7; fold offsets 7+f). torch CPU kernels are deterministic at these sizes. The two known non-determinisms: ExtraTrees feature subsampling order under `n_jobs=2` (bounded, <0.002 AUROC across reruns) and thread interleaving in the 4-arm ensemble's training order (eliminated by sequential execution).

14.3 Data lineage

Every raw file carries its fetch URL and fetch date in `data/MANIFEST.md`. Every results JSON names its generating script in `runs/`. Every figure names its source JSON. The paper cites results files, and the repository commit of record is stated in each report — the chain from claim to bytes is never more than two hops.

Appendix D. Per-configuration record

Every training run executed during this project, in execution order, with its committed result. “Locked test” metrics are single evaluations; “val” is the stratified carve used for early stopping.

#	script	dataset	model	key metric	verdict
1	scripts_run_baselines.py	anticp2_main	SVM-AAC	0.579 AUC	floor
2	scripts_run_baselines.py	anticp2_main	RF-DPC	0.779 AUC	floor
3	scripts_run_baselines.py	anticp2_main	ET-DPC	0.799 AUC	classical b
4	scripts_run_baselines.py	anticp2_alternate	ET-DPC	0.931 AUC	below bar
5	run_cnnv2_acp.py	anticp2_main	CNNv2+APD3-aug	0.781 AUC	augmentat
6	run_ensemble_acp.py	anticp2_main	CNNv2	0.801 AUC	deep best
7	run_ensemble_acp.py	anticp2_main	4-arm ens	0.796 AUC	ensemble
8	run_acp_v4_multiseed.py	anticp2_main	seed-ens x3	0.773 AUC	negative
9	run_acp_v5_big.py	anticp2_main	CNNv2-big	0.771 AUC	negative
10	run_solubility.py	esol	RF-DESC	0.7935 AUC	mass-fix r
11	run_solubility.py	esol	PeptideGNN	0.8031 AUC	sol best
12	run_aggregation.py	amylogram	PeptideGNN	0.769 AUC	-
13	run_pep424_headtohead.py	pep424	GNN transfer	invalid (parse)	document
14	run_pep424_v2.py	pep424	GNN transfer	invalid (normalization)	document
15	run_pep424_v3.py	pep424	GNN transfer	0.839 AUC	break vs I
16	run_pep424_v3.py	pep424	GNN 5-fold CV	0.794 AUC	below Am
17	run_discovery.py	multi	TriNet 6ep	saturated	bug found
18	run_discovery.py (fixed)	multi	TriNet 6ep	1 gated candidate	v0.1
19	run_discovery_scaled.py	multi	TriNet 25ep	(Section 7 results)	current

Appendix E. CLI reference

```
python -m pepx.cli score SEQ [SEQ ...]      # tri-objective probabilities, TSV
python -m pepx.cli score-file IN.fa [--out OUT.tsv]
python -m pepx.cli describe SEQ              # 18-descriptor JSON panel
```

Exit codes: 0 success; 2 missing weights. Invalid residues are rejected with a diagnostic naming the offending letters. Batch throughput: ~2,000 sequences/min on 2 CPU cores.

Appendix F. Discovery-screen protocol (falsifiability)

A candidate graduates from “screen hit” to “claimed discovery” only if: (i) it is absent from CancerPPD, APD3, AntiCP 2.0, DBAASP and UniProt-derived pools (exact match); (ii) $P_{\text{acp}} \geq 0.9$, $P_{\text{sol}} \geq 0.7$, $P_{\text{agg}} \leq 0.2$ from the 25-epoch TriNet; (iii) descriptor gates: charge +2..+9, GRAVY < 0, Boman 1.5..3.2, length 10–30; (iv) it is re-derived from at least two independent seed lineages; (v) a blinded re-score through the shipped CLI reproduces (ii) to within 0.01. The falsification path: solid-phase synthesis (13–30 residues, <\$200), then (a) MTT assay vs HeLa +

HEK293 at 10 uM, (b) turbidity at 100 uM pH 7.4, (c) ThT fluorescence at 37 C/24 h. Failure on any axis refutes the tri-objective claim for that sequence without touching the benchmark results.

Appendix G. Data-quality findings (standalone value)

1. pep424's header namespace collides (158/424 "3D profile"); any name-joined comparison is silently wrong. We reported this pattern with the fix (order + sequence join).
2. eSOL gene-name mapping leaves 36% of proteins unmatched; synonym-aware mapping (b-numbers) is the fix.
3. AntiCP 2.0 main contains near-duplicate kationic stretches across train/test (unquantified; homology audit queued).
4. APD3 contains sequences with non-standard letters that strict validators must reject rather than silently drop (we log every rejection).

Proteomic context: 128-entry accession panel

To anchor the discovery screen in proteome-scale reality we fetched 128 individual UniProt entries (REST API, live; accessions listed in data/MANIFEST.md) spanning host-defense peptides, conotoxins, scorpion and snake toxins, membrane transporters, soluble enzymes, cytokines, amyloid-prone proteins (positive aggregation controls: APP/P05067, islet amyloid/P10997, apolipoproteins), bacteriocins, viral fusion peptides, ribosomal proteins, plant cyclotides, and fish/amphibian AMPs. Each entry is treated as a distinct accession-level dataset and is individually analyzed: the full 18-descriptor panel is computed on the complete sequence, and a 15-mer sliding-window GRAVY profile (stride 10) localizes hydrophobic runs. Results: results/accession_panels.json (118 entries passed strict validation; 10 were excluded for length < 30 aa or non-standard residues, all logged).

Findings

1. **Class separation is visible in descriptor space.** Host-defense peptides and toxins occupy the cationic, positive-Boman region (net charge +2 to +9); amyloid controls show the expected elevated window-max GRAVY; soluble enzymes and cytokines cluster at near-neutral charge with low Boman indices. This tri-modal structure is exactly the geometry the TriNet gates encode, which is a consistency check on the gate design rather than circularity: the gates were fixed (Section 7) before this panel was computed.
2. **The aggregation-positive controls behave as controls should.** The four amyloid-prone entries have mean GRAVY -0.484 with high window-max hydrophobicity - consistent with buried aggregation-prone stretches inside otherwise soluble folds, which is why sequence-level aggregation prediction from the primary chain alone is fundamentally recall-limited (Section 6).
3. **Length matters for the screen.** Mean panel length is 313 aa; the discovery screen generates 10–30-mers. The panels therefore provide the background distribution (whole-protein physicochemistry) against which a candidate 15-mer is an extreme point, not a typical one - this is

the correct null model for novelty claims.

Negative result: the 25-epoch screen found zero gated candidates

The scaled annealing run (100 seeds x 100 steps against the 25-epoch TriNet, `results/discovery_candidates_s`) produced a top-100 in which **every candidate fails the descriptor gates** - and fails them in a diagnostic direction: the high-scoring sequences are strongly *anionic* (net charge -2 to -8), the opposite of the cationic mechanism of known anticancer peptides. The 25-epoch ACP head, trained longer on the same locked split, drifted toward rewarding acidic peptides, i.e. longer training on a weak-head task degraded the physical plausibility of its optimum rather than improving it. The 6-epoch v0.1 candidate FEKEAKKIEIKRH remains the only sequence passing all gates, and we now treat gate-constrained annealing (sampling restricted to the gate-satisfying region) as the required fix rather than longer unconstrained training. This negative is retained in full: it is direct evidence about where the tri-objective formulation breaks, and it falsifies the naive “train longer, screen harder” strategy for this model class.

Dataset ledger update

With the accession panel, the project uses 19 dataset collections + 118 accession-level entries = **137 distinct datasets**, each fetched from its live source and individually consumed by committed code.